

Machine Learning I

Ricardo Figueiredo up202105430

Sérgio Cardoso up202107918

May 14, 2025

Executive Summary

- **Goals:** use a classification algorithm taught in classes to obtain the accuracy of a set of benchmark datasets
- **Outline of the Approach:** modify an already made implementation of the algorithm on existing open-source code in order to make it robust to the data characteristic in question.
- **Summary of Results:**
 - Table comparing Accuracy/F1/Precision/Recall/AUC across datasets
 - Learning curves showing stability

Selected Algorithm

We chose the CART algorithm due to its easy learning curve. CART stands for Classification And Regression Tree which, in itself, uses a Decision Tree to determine the accuracy of a dataset. Decision Trees are among the most popular models in ML, as they are used not only by CART, but also by other algorithms, such as ID3 (Interactive Dichotomiser 3) or C4.5, the latter being an extension of CART.

CART uses the Gini index as its impurity measure:

Given a data set $D = \{\langle x_i, y_i \rangle\}_{i=1}^N$ (with N examples),

$Gini(D) = 1 - \sum_{i=1}^c p_i^2$, where p_i is the probability of class i .

Data Characteristic

For this assignment, we were given the three following set of benchmark datasets:

- 1 Noise or Outliers
- 2 Class Imbalance in Binary Classification
- 3 Multiclass Classification

Since CART uses binary partitioning, we ended up using the set of binary classification datasets (Class Imbalance in Binary Classification). We also used this set because of the fact that CART doesn't do well with class imbalanced data.

Experimental Setup:

- Datasets: binary classification datasets that suffer from class imbalance.
- Hyperparameters of the Algorithm:
 - Maximum Depth of the Tree
 - Criterion (Gini or Entropy)
 - Class Weight (Is it balanced or not?)

Performance Estimation Methodology:

Since most of the datasets tested for this assignment are so large, we used the holdout method, in which we split the data in two sub-sets: one for training and the other for testing.

Table comparing F1/AUC across datasets.
Learning curves showing stability.

Conclusions

References